

HBase

Theory and Practice of a Distributed Data Store

Pietro Michiardi

Eurecom

Introduction

Why yet another storage architecture?

- **Relational Database Management Systems (RDBMS):**

- ▶ Around since 1970s
- ▶ Countless examples in which they actually do make sense

- **The dawn of Big Data:**

- ▶ Previously: ignore data sources because no cost-effective way to store everything
 - ★ One option was to prune, by retaining only data for the last N days
- ▶ Today: store everything!
 - ★ Pruning fails in providing a base to build useful mathematical models

Batch processing

- **Hadoop and MapReduce:**

- ▶ Excels at storing (semi- and/or un-) structured data
- ▶ Data interpretation takes place at analysis-time
- ▶ Flexibility in data classification

- **Batch processing: A complement to RDBMS**

- ▶ Scalable sink for data, processing launched when time is right
- ▶ Optimized for large file storage
- ▶ Optimized for “streaming” access

- **Random Access:**

- ▶ Users need to “interact” with data, especially that “crunched” after a MapReduce job
- ▶ This is historically where RDBMS excel: random access for structured data

Column-Oriented Databases

- **Data layout:**

- ▶ Save their data grouped by columns
- ▶ Subsequent column values are stored contiguously on disk
- ▶ This is substantially different from traditional RDBMS, which save and store data by row

- **Specialized databases for specific workloads:**

- ▶ Reduced I/O
- ▶ Better suited for compression → Efficient use of bandwidth
 - ★ Indeed, column values are often very similar and differ little row-by-row
- ▶ Real-time access to data

- **Important NOTE:**

- ▶ HBase is not a column-oriented DB in the typical term
- ▶ HBase uses an on-disk column storage format
- ▶ Provides key-based access to specific cell of data, or a sequential range of cells

Column-Oriented and Row-Oriented storage layouts

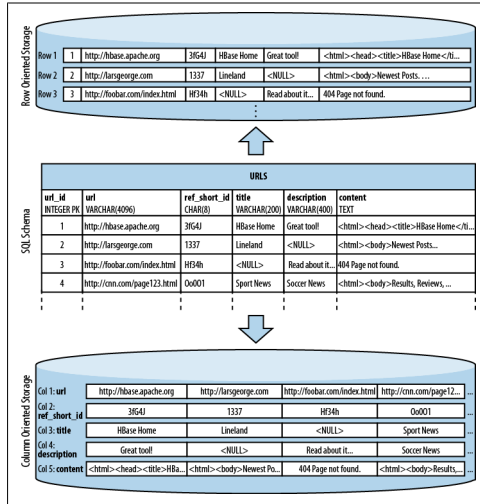


Figure: Example of Storage Layouts

The Problem with RDBMS

- **RDBMS are still relevant**

- ▶ Persistence layer for frontend application
- ▶ Store relational data
- ▶ Works well for a limited number of records

- **Example: Hush**

- ▶ Used throughout this course
- ▶ URL shortener service

- **Let's see the “scalability story” of such a service**

- ▶ Assumption: service must run with a reasonable budget

The Problem with RDBMS

- **Few thousands users: use a LAMP stack**

- ▶ *Normalize data*
- ▶ Use foreign keys
- ▶ Use Indexes

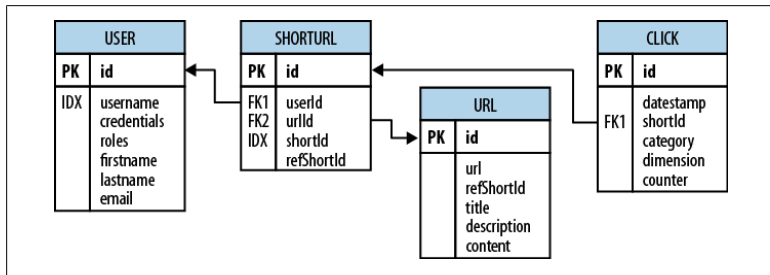


Figure: The Hush Schema expressed as an ERD

The Problem with RDBMS

- **Find all short URLs for a given user**

- ▶ JOIN `user` and `shorturl` tables
- ▶ Use the `WHERE` clause to select the given user

- **Stored Procedures**

- ▶ Consistently update data from multiple clients
- ▶ Underlying DB system guarantees coherency

- **Transactions**

- ▶ Make sure you can update tables in an *atomic* fashion
- ▶ RDBMS → *Strong Consistency* (ACID properties)
- ▶ *Referential Integrity*

The Problem with RDBMS

- **Scaling up to tens of thousands of users**

- ▶ Increasing pressure on the database server
- ▶ Adding more application servers is easy: they share their state on the same central DB
- ▶ CPU and I/O start to be a problem on the DB

- **Master-Slave architecture**

- ▶ Add DB server so that `READS` can be served in parallel
- ▶ Master DB takes all the writes (which are fewer in the Hush application)
- ▶ Slaves DB replicate Master DB and serve all reads (but you need a load balancer)

The Problem with RDBMS

● Scaling up to hundreds of thousands

- ▶ `READS` are still the bottlenecks
- ▶ Slave servers begin to fall short in serving clients requests

● Caching

- ▶ Add a caching layer, e.g. Memcached or Redis
- ▶ Offload `READS` to a fast in-memory system
- You lose consistency guarantees
- Cache invalidation is critical for having DB and Caching layer consistent

The Problem with RDBMS

● Scaling up more

- ▶ `WRITES` are the bottleneck
- ▶ The master DB is hit too hard by `WRITE` load
- ▶ *Vertical scalability*: beef up your master server
- This becomes costly, as you may also have to replace your RDBMS

● SQL JOINS becomes a bottleneck

- ▶ Schema de-normalization
- ▶ Cease using stored procedures, as they become slow and eat up a lot of server CPU
- ▶ Materialized views (they speed up `READS`)
- ▶ Drop secondary indexes as they slow down `WRITES`

The Problem with RDBMS

- **What if your application needs to further scale up?**

- ▶ Vertical scalability vs. Horizontal scalability

- **Sharding**

- ▶ Partition your data across multiple databases
 - ★ Essentially you break horizontally your tables and ship them to different servers
 - ★ This is done using fixed boundaries
 - Re-sharding to achieve load-balancing
- This is an operational nightmare
 - ▶ Re-sharding takes a huge toll on I/O resources

Non-Relational DataBases

- **They originally do not support SQL**

- ▶ In practice, this is becoming a thin line to make the distinction
- ▶ One difference is in the data model
- ▶ Another difference is in the consistency model (ACID and transactions are generally sacrificed)

- **Consistency models and the CAP Theorem**

- ▶ Strict: all changes to data are atomic
- ▶ Sequential: changes to data are seen in the same order as they were applied
- ▶ Causal: causally related changes are seen in the same order
- ▶ Eventual: updates propagates through the system and replicas when in steady state
- ▶ Weak: no guarantee

Dimensions to classify NoSQL DBs

• Data model

- ▶ How the data is stored: key/value, semi-structured, column-oriented, ...
- ▶ How to access data?
- ▶ Can the schema evolve over time?

• Storage model

- ▶ In-memory or persistent?
- ▶ How does this affect your access pattern?

• Consistency model

- ▶ Strict or eventual?
- ▶ This translates in how fast the system handles READS and WRITES [2]

Dimensions to classify NoSQL DBs

- **Physical Model**

- ▶ Distributed or single machine?
- ▶ How does the system scale?

- **Read/Write performance**

- ▶ Top-down approach: understands well the workload!
- ▶ Some systems are better for `READS`, other for `WRITES`

- **Secondary indexes**

- ▶ Does your workload require them?
- ▶ Can your system emulate them?

Dimensions to classify NoSQL DBs

● Failure Handling

- ▶ How each data store handle server failures?
- ▶ Is it able to continue operating in case of failures?
 - ★ This is related to Consistency models and the CAP theorem
- ▶ Does the system support “hot-swap”?

● Compression

- ▶ Is the compression method pluggable?
- ▶ What type of compression?

● Load Balancing

- ▶ Can the storage system seamlessly balance load?

Dimensions to classify NoSQL DBs

- **Atomic read-modify-write**

- ▶ Easy in a centralized system, difficult in a distributed one
- ▶ Prevent race conditions in multi-threaded or shared-nothing designs
- ▶ Can reduce client-side complexity

- **Locking, waits and deadlocks**

- ▶ Support for multiple client accessing data simultaneously
- ▶ Is locking available?
- ▶ Is it wait-free, hence deadlock free?

Impedance Match

“One-size-fits-all” has been long dismissed: need to find the perfect match for your problem.

Database (De-)Normalization

- **Schema design at scale**

- ▶ A good methodology is to apply the DDI principle [8]
 - ★ Denormalization
 - ★ Duplication
 - ★ Intelligent Key design

- **Denormalization**

- ▶ Duplicate data in more than one table such that at `READ` time no further aggregation is required

- **Next: an example based on Hush**

- ▶ How to convert a classic relational data model to one that fits HBase
- ▶ This example will be covered in the LAB session 3

Example: Hush - from RDBMS to HBase




Figure: The Hush Schema expressed as an ERD

- `shorturl` table: contains the short URL
- `click` table: contains click tracking, and other statistics, aggregated on a daily basis (essentially, a counter)
- `user` table: contains user information
- `URL` table: contains a replica of the page linked to a short URL, including META data and content (this is done for batch analysis purposes)

Example: Hush - from RDBMS to HBase




Figure: The Hush Schema expressed as an ERD

- `user` table is indexed on the `username` field, for fast user lookup
- `shorturl` table is indexed on the short URL (`shortId`) field, for fast short URL lookup

Example: Hush - from RDBMS to HBase




Figure: The Hush Schema expressed as an ERD

- `shorturl` and `user` tables are related through a foreign key relation on the `userId`
- `URL` table is related to `shorturl` table with a foreign key on the `URL id`
- `click` table is related to `shorturl` table with a foreign key on the `short URL id`
- **NOTE:** a web page is stored only once (even if multiple users link to it), but each users maintain separate statistics

Example: Hush - from RDBMS to HBase

Table: shorturl		
Row Key:	shortId	
Family:	data:	Columns: url, refShortId, userId, clicks
	stats-daily: [ttl: 7days]	Columns: YYYYMMDD, YYYYMMDD\x00<country-code>
	stats-weekly: [ttl: 4weeks]	Columns: YYYYWW, YYYYWW\x00<country-code>
	stats-monthly: [ttl: 12months]	Columns: YYYYMM, YYYYMM\x00<country-code>

Table: url		
Row Key:	MD5(url)	
Family:	data: [compressed]	Columns: refShortId, title, description
	content: [compressed]	Columns: raw

Table: user-shorturl		
Row Key:	username\x00shortId	
Family:	data:	Columns: timestamp

Table: user		
Row Key:	username	
Family:	data:	Columns: credentials, roles, firstname, lastname, email

- `shorturl` table: stores each short URL, usage statistics (various time-ranges in separate *column-families* with distinct *TTL* settings)
 - ▶ Note the dimensional postfix appended to the time information
- `url` table: stores the downloaded page, and the extracted details
 - ▶ This table uses compression

Figure: The Hush Schema in HBase

Example: Hush - from RDBMS to HBase

Table: shorturl		
Row Key:	shortId	
Family:	data:	Columns: url, refShortId, userId, clicks
	stats-daily: [ttl: 7days]	Columns: YYYYMMDD, YYYYMMDD\x00<country-code>
	stats-weekly: [ttl: 4weeks]	Columns: YYYYWW, YYYYWW\x00<country-code>
	stats-monthly: [ttl: 12months]	Columns: YYYYMM, YYYYMM\x00<country-code>

Table: url		
Row Key:	MD5(url)	
Family:	data: [compressed]	Columns: refShortId, title, description
	content: [compressed]	Columns: raw

Table: user-shorturl		
Row Key:	username\x00shortId	
Family:	data:	Columns: timestamp

Table: user		
Row Key:	username	
Family:	data:	Columns: credentials, roles, firstname, lastname, email

- `user-shorturl` table: this is a lookup table (basically an index) to find all shortIDs for a given user

- ▶ Note that this table is filled at *insert time*, it's not automatically generated by HBase

- `user` table: stores user details

Figure: The Hush Schema in HBase

Example: Hush - RDBMS vs HBase

- **Same number of tables**

- ▶ Their meaning is different
- ▶ `click` table has been absorbed by the `shorturl` table
- ▶ statistics are stored with the date as the key, so that they can be accessed *sequentially*
- ▶ The `user-shorturl` table is replacing the foreign key relationship, making user-related lookups faster

- **Normalized vs. De-normalized data**

- ▶ Wide tables and column-oriented design eliminates `JOINS`
- ▶ *Compound keys* are essential
- ▶ Data partitioning is based on keys, so a proper understanding thereof is essential

HBase building blocks

- **The backdrop: BigTable**

- ▶ GFS, The Google FileSystem [6]
- ▶ Google MapReduce [4]
- ▶ BigTable [3]

- **What is BigTable?**

- ▶ BigTable is a distributed storage system for managing structured data designed to scale to a very large size
- ▶ BigTable is a sparse, distributed, persistent multi-dimensional sorted map

- **What is HBase?**

- ▶ Essentially it's an open-source version of BigTable
- ▶ Differences listed in [5]

HBase building blocks

Tables, Rows, Columns, and Cells

- **The most basic unit in HBase is a *column***

- ▶ Each column may have multiple versions, with each distinct value contained in a separate *cell*
- ▶ One or more columns form a *row*, that is addressed uniquely by a *row key*

- A table is a collection of rows

- ▶ All rows are always *sorted lexicographically* by their row key

```
hbase(main):001:0> scan 'table1'
ROW                                COLUMN+CELL
row-1                             column=cf1:, timestamp=1297073325971 ...
row-10                            column=cf1:, timestamp=1297073337383 ...
row-11                            column=cf1:, timestamp=1297073340493 ...
row-2                             column=cf1:, timestamp=1297073329851 ...
row-22                            column=cf1:, timestamp=1297073344482 ...
row-3                             column=cf1:, timestamp=1297073333504 ...
row-abc                           column=cf1:, timestamp=1297073349875 ...
7 row(s) in 0.1100 seconds
```

HBase building blocks

Tables, Rows, Columns, and Cells

- **Lexicographical ordering of row keys**

- ▶ Keys are compared on a binary level, byte by byte, from left to right
- ▶ This can be thought of as a primary index on the row key!
- ▶ Row keys are *always unique*
- ▶ Row keys can be any *arbitrary array of bytes*

- **Columns**

- ▶ Rows are composed of columns
- ▶ Can have millions of columns
- ▶ Can be compressed or tagged to stay in memory

HBase building blocks

Tables, Rows, Columns, and Cells

● Column Families

- ▶ Columns are grouped into *column families*
- Semantical boundaries between data
- ▶ Column families and columns stored together in the same low-level storage file, called an *HFile*
- ▶ Defined when table is created
- ▶ Should not be changed too often
- ▶ The number of column families should be reasonable [WHY?]
- ▶ Column family name composed by printable characters

● References to columns

- ▶ Column “name” is called *qualifier*, and can be any arbitrary number of bytes
- ▶ Reference: `family:qualifier` (also called the **column key**)

HBase building blocks

Tables, Rows, Columns, and Cells

- **A note on the `NULL` value**

- ▶ In RDBMS `NULL` cells need to be set and occupy space
- ▶ In HBase, `NULL` cells or columns are simply not stored

- **A *cell***

- ▶ Every column value, or cell, is timestamped (implicitly or explicitly)
 - ★ This can be used to save multiple versions of a value that changes over time
 - ★ Versions are stored in decreasing timestamp, most recent first
- ▶ Cell versions can be constrained by *predicate deletions*
 - ★ Keep only values from the last week

HBase building blocks

Tables, Rows, Columns, and Cells

● Access to data

- ▶ (Table, RowKey, Family, Column, Timestamp) → Value
- ▶ `SortedMap<RowKey, List<SortedMap<Column, List<Value, Timestamp>>>>`
- ▶ The first `SortedMap` is the table, containing a `List` of column families
- ▶ The families contain another `SortedMap`, representing columns and a `List` of value, timestamp tuples

● A note on consistency:

- ▶ Row data access is **atomic** and includes any number of columns
 - ▶ There is no further guarantee or transactional feature spanning multiple rows
- HBase is strictly consistent

HBase building blocks

Automatic Sharding

● Region

- ▶ This is the basic unit of scalability and load balancing
- ▶ Regions are contiguous ranges of rows “stored together” → they are the equivalent of *range partitions* in sharded RDBMS
- ▶ Regions are *dynamically split* by the system when they become too large
- ▶ Regions can also be merged to reduce the number of storage files

● Regions in practice

- ▶ Initially, there is one region
- ▶ System monitors region size: if a threshold is attained, `SPLIT`
 - ★ Regions are split in two at the *middle key*
 - ★ This creates roughly two equivalent (in size) regions

HBase building blocks

Automatic Sharding

- **Region Servers**

- ▶ Each region is served by *exactly one Region Server*
- ▶ Region servers can serve multiple regions
- ▶ The number of region servers and their sizes depend on the capability of a single region server

- **Server failures**

- ▶ Regions allow for fast recovery upon failure
- ▶ Fine-grained Load Balancing is also achieved using regions as they can be easily moved across servers

HBase building blocks

Storage API

- **No support for SQL**

- ▶ CRUD operations using a standard API, available for many “clients”
- ▶ Data access is not declarative but imperative

- **Scan API**

- ▶ Allows for fast iteration over ranges of rows
- ▶ Allows to limit the number and which column are returned
- ▶ Allows to control the version number of each cell

- **Read-modify-write API**

- ▶ HBase supports single-row transactions
- ▶ Atomic read-modify-write on data stored in a single row key

HBase building blocks

Storage API

● Counters

- ▶ Values can be interpreted as counters and **updated atomically**
- ▶ Can be read and modified in one operation
- Implement global, strictly consistent, sequential counters

● Coprocessors

- ▶ These are equivalent to stored-procedures in RDBMS
- ▶ Allow to push user code in the address space of the server
- ▶ Access to server local data
- ▶ Implement lightweight batch jobs, data pre-processing, data summarization

HBase building blocks

HBase implementation

• Data Storage

- ▶ Store files are called `HFiles`
- ▶ Persistent and ordered **immutable** maps from key to value
- ▶ Internally implemented as sequences of blocks with an index at the end
- ▶ Index is loaded when the `HFile` is opened and kept in memory

• Data lookups

- ▶ Since `HFiles` have a block index, lookup can be done with a single disk seek
- ▶ First, the block possibly containing a given lookup key is determined with a **binary search** in the in-memory index
- ▶ Then a block read is performed to find the actual key

• Underlying file system

- ▶ Many are supported, usually HBase deployed on top of HDFS

HBase building blocks

HBase implementation

● **WRITE operation**

- ▶ First, data is written to a commit log, called WAL (write-ahead-log)
- ▶ Then data is moved into memory, in a structure called `memstore`
- ▶ When the size of the `memstore` exceeds a given threshold it is flushed to an `HFile` to disk

● **How can HBase write, while serving READS and WRITES?**

- ▶ Rolling mechanism
 - ★ new/empty slots in the `memstore` take the updates
 - ★ old/full slots are flushed to disk
- ▶ Note that data in `memstore` is sorted by keys, matching what happens in the `HFiles`

● **Data Locality**

- ▶ Achieved by the system looking up for server hostnames
- ▶ Achieved through intelligent key design

HBase building blocks

HBase implementation

• Deleting data

- ▶ Since HFiles are immutable, how can we delete data?
- ▶ A delete marker (also known as *tombstone marker*) is written to indicate that a given key is deleted
- ▶ During the read process, data marked as deleted is skipped
- ▶ Compactions (see next slides) finalize the deletion process

• READ operation

- ▶ Merge of what is stored in the `memstores` (data that is not on disk) and in the `HFiles`
- ▶ The WAL is never used in the `READ` operation
- ▶ Several API calls to read, scan data

HBase building blocks

HBase implementation

● Compactions

- ▶ Flushing data from `memstores` to disk implies the creation of new `HFiles` each time
- We end up with many (possibly small) files
- We need to do housekeeping [**WHY?**]

● Minor Compaction


- ▶ Rewrites small `HFiles` into fewer, larger `HFiles`
- ▶ This is done using an n -way merge¹

● Major Compaction

- ▶ Rewrites all files within a column family or a region in a new one
- ▶ Drop deleted data
- ▶ Perform predicated deletion (e.g. delete old data)

¹What is MergeSort?

HBase: a glance at the architecture



- **Master node: HMaster**

- ▶ Assigns regions to region servers using ZooKeeper
- ▶ Handles load balancing
- ▶ Not part of the data path
- ▶ Holds metadata and schema

- **Region Servers**

- ▶ Handle `READs` and `WRITEs`
- ▶ Handle region splitting

Architecture

Seek vs. Transfer

- **Fundamental difference between RDBMS and alternatives**

- ▶ B+Trees
- ▶ Log-Structured Merge Trees

- **Seek vs. Transfer**

- ▶ Random access to individual cells
- ▶ Sequential access to data

B+ Trees

- **Dynamic, multi-level indexes**

- ▶ Efficient insertion, lookup and deletion
- ▶ **Q: What's the difference between a B+ Tree and a Hash Table?**
- ▶ Frequent updates may imbalance the trees → Tree optimization and re-organization is required (which is a costly operation)

- **Bounds on *page size***

- ▶ Number of keys in each branch
- ▶ Larger fanout compared to binary trees
- ▶ Lower number of I/O operations to find a specific key

- **Support for range scans**

- ▶ Leafs are linked and represent an in-order list of all keys
- ▶ No costly tree-traversal algorithms required

LSM-Trees

• Data flow

- ▶ Incoming data is first stored in a logfile, *sequentially*
- ▶ Once the log has the modification saved, data is pushed in memory
 - ★ In-memory store holds most recent updates for fast lookup
- ▶ When memory is “full”, data is flushed in a store file to disk, as a sorted list of `key` \rightarrow `record` pair
- ▶ At this point, the log file can be thrown away

• How store files are arranged

- ▶ Similar idea of a B+ Tree, but optimized for sequential disk access
- ▶ All nodes of the tree try to be filled up completely
- ▶ Updates are done in a **rolling merge** fashion
 - ★ The system packs existing on-disk multi-page blocks with in-memory data until the block reaches full capacity

LSM-Trees

● Clean-up process

- ▶ As flushes take place over time, a lot of store files are created
- ▶ Background process aggregates files into larger ones to limit disk seeks
- ▶ All store files are always sorted by key → no re-ordering required to fit new keys in

● Data Lookup

- ▶ Lookups are done in a merging fashion
 - ★ First lookup in the in-memory store
 - ★ If miss, the lookup in the on-disk store

● Deleting data

- ▶ Use a *delete marker*
- ▶ When pages are re-written, deleted markers and keys are eventually dropped
- ▶ Predicate deletion happens here

B+ Tree vs. LSM-Trees

● B+ Tree [1]

- ▶ Work well when there are not so many updates
- ▶ The more and the faster you insert data at random locations the faster pages get fragmented
- ▶ **Updates and deletes are done at disk seek rates, rather than transfer rates**

● LSM-Tree [7]

- ▶ Work at disk transfer rate and scale better to huge amounts of data
- ▶ Guarantee a consistent insert rate
 - ★ They transform random into sequential writes
- ▶ Reads are independent from writes
- ▶ Optimized data layout which offers predictable boundaries on disk seeks

Overview




Figure: Overview of how HBase handles files in the filesystem

Storage

Overview

- **HBase handles two kinds of file types**

- ▶ One is used for the WAL
- ▶ One is used for the actual data storage

- **Who does what**

- ▶ `HMaster`
 - ★ Low-level operations
 - ★ Assigns region servers to key space
 - ★ Keeps metadata
 - ★ Talks to ZooKeeper
- ▶ `HRegionServer`
 - ★ Handles the WAL and `HFiles`
 - ★ These files are divided in to blocks and stored into HDFS
 - ★ Block size is a parameter

Storage

Overview

● General communication flow

- ▶ A client contacts ZooKeeper when trying to access a particular row
- ▶ Recovers from ZooKeeper the server name that host the `-ROOT-` region
- ▶ Using the `-ROOT-` information the client retrieves the server name that host the `.META.` table region
 - ★ The `.META.` table region contains the row key in question
- ▶ Contact the reported `.META.` server and retrieve the server name that has the region containing the row key in question

● Caching

- ▶ Generally, lookup procedures involve caching row key locations for faster subsequent lookups

Storage

Overview

● Important Java Classes

- ▶ `HRegionServer` handles one or more regions and create the corresponding `HRegion` object
- ▶ When an `HRegion` object is opened it creates a `Store` instance for each `HColumnFamily`
- ▶ Each `Store` instance can have:
 - ★ One or more `StoreFile` instances
 - ★ A `MemStore` instance
- ▶ `HRegionServer` has a shared `HLog` instance

Storage

Write Path

- **External client insert data in HBase**

- ▶ Issues an `HTable.put(Put)` request to `HRegionServer`
- ▶ `HRegionServer` hands the request to the `HRegion` instance that matches the request [Q: What is the matching criteria?]

- **How the system reacts to a write request**

- ▶ Write data to the WAL, represented by the `HLog` class
 - ★ The WAL stores `HLogKey` instances in a HDFS `SequenceFile`
 - ★ These keys contain a sequence number and the actual data
 - ★ In case of failure, this data can be used to replay not-yet-persisted data
- ▶ Copy data in the `MemStore`
 - ★ Check if `MemStore` size has reached a threshold
 - ★ If yes, launch a *flush request*
 - ★ Launch a thread in the `HRegionServer` and flush `MemStore` data to an `HFile`

Storage

HBase Files

- **What and where are HBase files (including WAL, HFile,...) stored?**
 - ▶ HBase has a root directory set to “/hbase” in HDFS
 - ▶ Files can be divided into:
 - ★ Those that reside under the HBase root directory
 - ★ Those that are in the *per-table* directories
- /hbase
 - ▶ .logs
 - ▶ .oldlogs
 - ▶ .hbase.id
 - ▶ .hbase.version
 - ▶ /example-table

Storage

HBase Files

- /example-table
 - ▶ .tableinfo
 - ▶ .tmp
 - ▶ "...Key1..."
 - ★ .oldlogs
 - ★ .regioninfo
 - ★ .tmp
 - ★ colfam1/
- colfam1/
 - ▶ "...column-key1..."

Storage

HBase: Root-level files

- **.logs directory**

- ▶ WAL files handled by `HLog` instances
- ▶ Contains a subdir for each `HRegionServer`
- ▶ Each subdir contains many `HLog` files
- ▶ All regions from that `HRegionServer` share the same `HLog` files

- **.oldlogs directory**

- ▶ When data is persisted to disk (from `Memstores`) log files are decommissioned to the `.oldlogs` dir

- **hbase.id and hbase.version**

- ▶ Represent the unique ID of the cluster and the file format version

Storage

HBase: Table-level files

- **Every table has its own directory**

- ▶ `.tableinfo`: stores the serialized `HTableDescriptor`
 - ★ This include the table and column family schema
- ▶ `.tmp` directory
 - ★ Contains temporary data

Storage

HBase: Region-level files

- **Inside each table dir, there is a separate dir for every region in the table**
 - ▶ The name of each of this dirs is the MD5 hash of a region name
 - ★ Inside each region there is a directory for each column family
 - ★ Each column family directory holds the actual data files, namely `HFiles`
 - ★ Their name is just an arbitrary random number
 - ▶ Each region directory also has a `.regioninfo` file
 - ★ Contains the serialized information of the `HRegionInfo` instance
- **Split Files**
 - ▶ Once the region needs to be split, a `splits` directory is created
 - ★ This is used to stage two daughter regions
 - ★ If split is successful, daughter regions are moved up to the table directory

Storage

HBase: A note on region splits

- **Splits triggered by store file (region) size**
 - ▶ Region is split in two
 - ▶ Region is closed to new requests
 - ▶ `.META.` is updated
- **Daughter regions initially reside on the same server**
 - ▶ Both daughters are compacted
 - ▶ Parent is cleaned up
 - ▶ `.META.` is updated
- **Master schedules new regions to be moved off to other servers**

Storage

HBase: Compaction

- **Process that takes care of re-organizing store files**
 - ▶ Essentially to conform to underlying filesystem requirements
 - ▶ Compaction check when memstore is flushed
- **Minor and Major compactions**
 - ▶ Always from the oldest to the newest files
 - ▶ Avoid all servers to perform compaction concurrently




Figure: A set of store files showing the minimum compaction threshold

Storage

HFile format

- **Store files are implemented by the HFile class**
 - ▶ Efficient data storage is the goal
- **HFiles consist of a variable number of blocks**
 - ▶ Two fixed blocks: *info* and *trailer*
 - ▶ *index* block: records the offsets of the *data* and *meta* blocks
 - ▶ Block size: *large* → sequential access; *small* → random access




Figure: The HFile structure

Storage

HFile size and HDFS block size

- **HBase uses any underlying filesystem**
- **In case HDFS is used**
 - ▶ HDFS block size is generally 64MB
 - ▶ This is 1,024 times the default `HFile` block size (64 KB)
 - There is no correlation between HDFS block and HFile sizes

Storage

The KeyValue Format

- **Each KeyValue in the HFile is a low-level byte array**
 - ▶ It allows for *zero-copy* access to the data
- **Format**
 - ▶ Fixed-length preamble indicates the length of the key and value
 - ★ This is useful to offset into the array to get direct access to the value, ignoring the key
 - ▶ Key format
 - ★ Contains row key, column family name, column qualifier...
 - ★ [TIP]: consider small keys to avoid overhead when storing small data




Figure: The KeyValue Format

The Write-Ahead Log

- **Main tool to ensure resiliency to failures**

- ▶ Region servers keep data in-memory until enough is collected to warrant a flush
- ▶ What if the server crashes or power is lost?

- **WAL is a common approach to address fault-tolerance**

- ▶ Every data update is first written to a log
- ▶ Log is persisted (and replicated, since it resides on HDFS)
- ▶ Only when log is written, client is notified a successful operation on data

The Write-Ahead Log

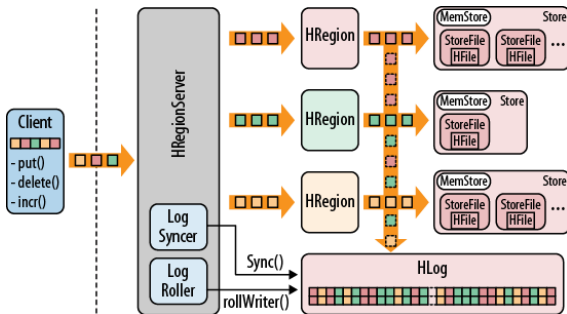



Figure: The write path of HBase

- **WAL records all changes to data**

- ▶ Can be replayed in case of server failure
- ▶ If write to WAL fails, the whole operations has to fail

The Write-Ahead Log



Write Path

- ▶ Client modifies data (`put()`, `delete()`, `increment()`)
- ▶ Modifications are wrapped into a **KeyValue** object
- ▶ Objects are batched to the corresponding **HRegionServer**
- ▶ Objects are routed to the corresponding **HRegion**
- ▶ Objects are written to **WAL** and in the **MemStore**

Read Path

- **HBase uses multiple store files per column family**
 - ▶ These can be either in-memory and/or materialized on disk
 - ▶ Compactions and clean-up background processes take care of store files maintenance
 - ▶ Store files are immutable, so deletion is handled in a special way
- **The anatomy of a get command**
 - ▶ HBase uses a `QueryMatcher` in combination with a `ColumnTracker`
 - ▶ First, an exclusion check is performed to filter skip files (and eventually tombstone labelled data)
 - ▶ Scanning data is implemented by a `RegionScanner` class which retrieves a `StoreScanner`
 - ▶ `StoreScanner` includes both the `MemStore` and `HFiles`
 - ▶ Read/Scans happen in the same order as data is saved

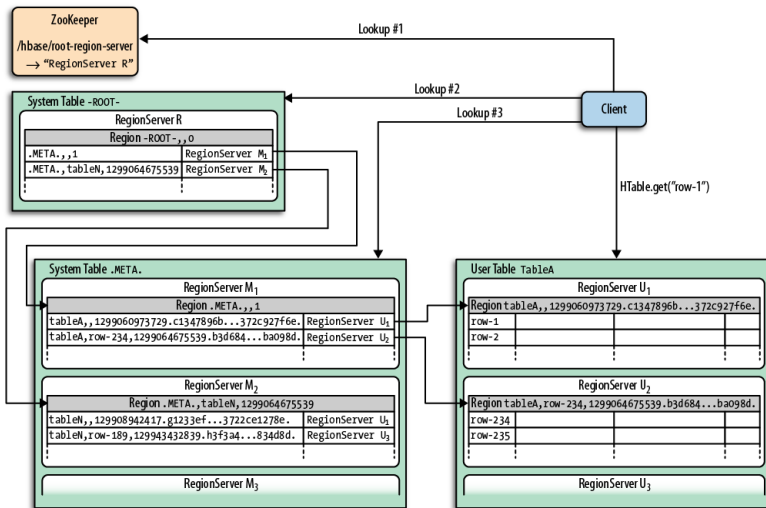
Region Lookups

- **How does a client find the region server hosting a specific row key range?**
 - ▶ HBase uses two special catalog tables, `-ROOT-` and `.META.`
 - ▶ The `-ROOT-` table is used to refer to all regions in the `.META.` table
- **Three-level B+ Tree -like operation**
 - ▶ Level 1: a node stored in ZooKeeper, containing the location (region server) of the `-ROOT-` table
 - ▶ Level 2: Lookup in the `-ROOT-` table to find a matching meta region
 - ▶ Level 3: Retrieve the table region from the `.META.` table

Region Lookups

- **Where to send requests when looking for a specific row key?**
 - ▶ This information is cached, but the first time or when the cache is stale or when there is a miss due to compaction, the following procedure applies
- **Recursive discovery process**
 - ▶ Ask the region server hosting the matching `.META.` table to retrieve the row key address
 - ▶ If the information is invalid, it backs out: asks the `-ROOT-` table where the relevant `.META.` region is
 - ▶ If this fails, ask ZooKeeper where the `-ROOT-` table is

Region Lookups



Key Design

Concepts

- **HBase has two fundamental key structures**

- ▶ Row key
- ▶ Column key

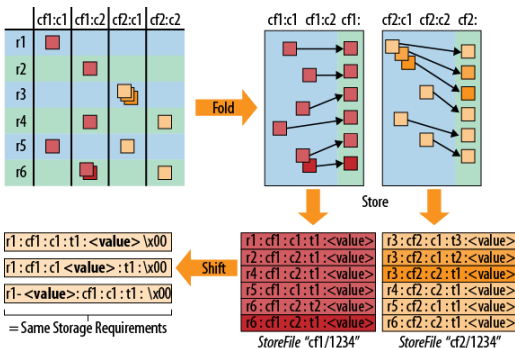
- **Both can be used to convey meaning**

- ▶ Because they store particularly meaningful data
- ▶ Because their sorting order is important


Concepts

● Logical vs. on-disk layout of a table

- ▶ Main unit of separation within a table is the *column family*
- ▶ The actual columns (as opposed to other column-oriented DB) are not used to separate data
- ▶ Although cells are stored logically in a table format, rows are stored as linear sets of the cells
- ▶ Cells contain all the vital information inside them




Concepts



Logical Layout (Top-Left)

- ▶ Table consists of rows and columns
 - ▶ Columns are the combination of a column family name and a column qualifier
- `<cf name: qualifier>` is the **column key**
- ▶ Rows have a **row key** to address all columns of a single logical row

Concepts



Folding the Logical Layout (Top-Right)

- ▶ The cells of each row are stored one after the other
- ▶ Each column family are stored separately
- On disk all cells of one family reside on an individual StoreFile
- ▶ HBase does not store unset cells
- **Row and column key is required to address every cell**

Concepts


● Versioning

- ▶ Multiple versions of the same cell stored consecutively, together with the *timestamp*
- ▶ Cells are sorted in descending order of timestamp
- Newest value first

● KeyValue object

- ▶ The entire cell, with all the structural information, is a KeyValue object
- ▶ Contains: row key, <column family: qualifier> → column key, timestamp and value
- ▶ Sorted by row key first, then by column key

Concepts



Physical Layout (Lower-Right)

- ▶ Select data by row key
 - ★ This reduces the amount of data to scan for a row or a range of rows
- ▶ Select data by row key and column key
 - ★ This focuses the system on an individual storage file
- ▶ Select data by column qualifier
 - ★ Exact lookups, including filters to omit useless data

Concepts

- Summary of key lookup properties

Key Value	Key				Value
	Row	Column Family	Column Qualifier	Timestamp	
Skip Rows	✓	X	X	X	X
Skip Store Files	✓	✓	X	✓	X
Filter Compatible	✓	✓	✓	✓	✓

← Performance

Increased Cardinality →

Tall-Narrow vs. Flat-Wide Tables

- **Tall-Narrow Tables**

- ▶ Few columns
- ▶ Many rows

- **Flat-Wide Tables**

- ▶ Many columns
- ▶ Few rows

- **Given the query granularity explained before**

- Store parts of the cell data in the row key
 - ▶ Furthermore, HBase splits at row boundaries
- It is recommended to go for Tall-Narrow Tables

Tall-Narrow vs. Flat-Wide Tables

● Example: email data - version 1

- ▶ You have all emails of a user in a single row (e.g. `userID` is the row key)
- ▶ There will be some outliers with orders of magnitude more emails than others
- A single row could outgrow the maximum file/region size and work against split facility

● Example: email data - version 2

- ▶ Each email of a user is stored in a separate row (e.g. `userID:messageID` is the row key)
- ▶ On disk this makes no difference (see the disk layout figure)
 - ★ If the `messageID` is in the column qualifier or the row key, each cell still contains a single email message
- The table can be split easily and the query granularity is more fine-grained

Partial Key Scans

- **Partial Key Scans reinforce the concept of Tall-Narrow Tables**

- ▶ From the email example: assume you have a separate row per message, across all users
- ▶ If you don't have an exact combination of user and message ID you cannot access a particular message

- **Partial Key Scan solves the problems**

- ▶ Specify a *start* and *end* key
- ▶ The start key is set to the exact `userID` only, with the end key set at `userID+1`
- This triggers the internal lexicographic comparison mechanism
 - ★ Since the table does not have an exact match, it positions the scan at:
`<userID>:<lowest-messageID>`
- ▶ The scan will then iterate over all the messages of an exact user, parse the row key and get the `messageID`

Partial Key Scans

- **Composite keys and atomicity**

- ▶ Following the email example: a single user inbox now spans many rows
- ▶ It is no longer possible to modify a single user inbox in one atomic operation

- **If this is acceptable or not, depends on the application at hand**

Time Series Data

- **Stream processing of events**

- ▶ E.g. data coming from a sensor, stock exchange, monitoring system ...
- ▶ Such data is a time series → **The row key represents the event time**
- HBase will store all rows sorted in a distinct range, namely regions with specific start and stop keys

- **Sequential monotonously increasing nature of time series data**

- ▶ All incoming data is written to the same region (and hence the same server)
- **Regions become HOT!**
- ▶ Performance of the whole cluster is bound to that of a single machine

Time Series Data

● Solution to achieve load balancing: Salting

- ▶ We want data to be spread over all region servers
- ▶ This can be done, e.g., by prefixing the row key with a non-sequential number

Salting example

```
byte prefix = (byte) (Long.hashCode(timestamp) % <number of  
region servers>);  
byte[] rowkey = Bytes.add(Bytes.toBytes(prefix),  
Bytes.toBytes(timestamp));
```

- Data access needs to be *fanned out* across many servers
- + Use multiple threads to read for I/O performance: e.g. use the Map phase of MapReduce

Time Series Data

- **Solution to achieve load balancing: Field swap/promotion**
 - ▶ Move the timestamp field of the row key or prefix it with another field
 - ★ If you already have a composite row key, simply *swap* elements
 - ★ Otherwise if you only have the timestamp, you need to *promote* another field
 - ▶ The sequential, monotonously increasing timestamp is moved to a secondary position in the row key
- You can only access data (especially time ranges) for a given swapped or promoted field (but this could be a feature)
- + You achieve load balancing

Time Series Data


- **Solution to achieve load balancing: Randomization**

- ▶ `byte[] rowkey = MD5(timestamp)`
- ▶ This gives you a random distribution of the row key across all available region servers

- Less than ideal for range scans
- + Since you can re-hash the timestamp, this solution is good for **random access**

Time Series Data

- Summary



MapReduce Integration

Introduction

- **In the following we review the main classes involved in reading and writing data from/to an underlying data store**
- **For MapReduce to work with HBase, some more practical issues have to be addressed**
 - ▶ E.g.: creating an appropriate JAR file inclusive of all required libraries
- Refer to [5], Chapter 7 for an in-depth treatment of this subject

Main classes involved in MapReduce

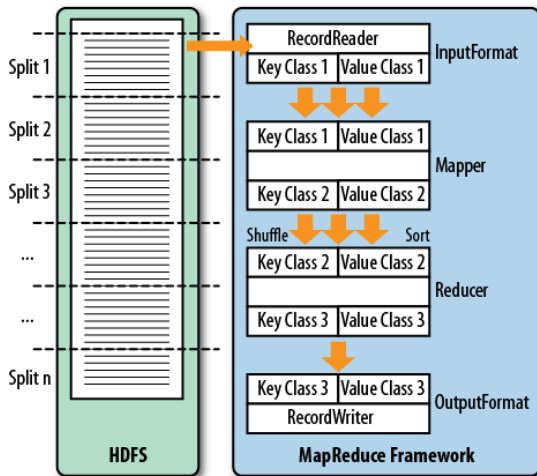


Figure: Main MapReduce Classes

Main classes involved in MapReduce

InputFormat

- **It is responsible for two things**

- ▶ Splits input data
- ▶ Returns a `RecordReader` instance
 - ★ Defines a *key* and a *value* object
 - ★ Provides a `next()` method to iterate over input records




Figure: InputFormat hierarchy

Main classes involved in MapReduce

`InputFormat` → `TableInputFormatBase`

- **Implement a full turnkey solution to scan an HBase table**
 - ▶ Splits the table into proper blocks and hand them to the MapReduce process
- **Must supply a Scan instance to interact with a table**
 - ▶ Specify start and stop keys for the scan
 - ▶ Add filters (optional)
 - ▶ Specify the number of versions

Main classes involved in MapReduce

Mapper

- Each record read using the `RecordReader` is processed using the `map()` method
- The `Mapper` reads specific types of input key/value pairs, but emit possibly another type



Figure: The Mapper hierarchy

Main classes involved in MapReduce

Mapper → TableMapper

- **TableMapper class enforces:**

- ▶ The input key to the mapper to be an `ImmutableBytesWritable` type
- ▶ The input value to be a `Result` type

- **A handy implementation is the `IdentityTableMapper`**

- ▶ This is the equivalent of an identity mapper

Main classes involved in MapReduce

OutputFormat

- **Used to persist data**

- ▶ Output written to files
- ▶ Output written to HBase tables
 - ★ This is done using a `TableRecordWriter`




Figure: The OutputFormat hierarchy

Main classes involved in MapReduce

OutputFormat → TableOutputFormat

- **This is the class that handles the key/value pairs and writes them to their final destination**
 - ▶ Single instance that takes the output record from each reducer subsequently
- **Details**
 - ▶ Must specify the table name when the MR job is created
 - ▶ Handles buffer flushing implicitly (*autoflush* option is set to false)

MapReduce Locality

- **How does the system make sure data is placed close to where it is needed?**

- ▶ This is done implicitly by MapReduce when using HDFS
- ▶ When MapReduce uses HBase things are a bit different

- **How HBase handles data locality**

- ▶ Shared vs. non-shared cluster
- ▶ HBase store its files on HDFS (`HFiles` and `WAL`)
- ▶ HBase servers are not restarted frequently and they perform compactions regularly
- HDFS is smart enough to ensure data locality
 - ★ There is a block placement policy that enforces local writes
 - ★ The data node compares the server name of the writer with its own
 - ★ If they match, the block is written to the local filesystem
- ▶ Just be careful about region movements during load balancing or server failures

Table Splits

- **When running a MapReduce job that reads from an HBase table you use the `TableInputFormat`**
 - ▶ Overrides `getSplits()` and `createRecordReader()`
- **Before a job is run, the framework calls `getSplit()` to determine how the data is to be separated into chunks**
 - ▶ `TableInputFormat`, given the `Scan` instance you define, divide the table at region boundaries
 - The number of input splits is equal to all regions between the start and stop keys

Table Splits

- **When a job starts, the framework calls `createRecordReader()` for each input split**
 - ▶ It iterates over the splits and create a new `TableRecordReader` with the current split
 - ▶ Each `TableRecordReader` handles exactly one region, reading and mapping every row from the region's start and end keys
- **Data locality**
 - ▶ Each split contains the server name hosting the region
 - ▶ The framework checks the server name and if the `TaskTracker` is running on the same machine, it will run it on that server
 - ▶ The `RegionServer` is colocated with the HDFS `DataNode`, hence data is read from the local filesystem
- **TIP: Turn off speculative execution!**

References I

[1] B+ tree.

http://en.wikipedia.org/wiki/B%2B_tree.

[2] Eric Brewer.

Lessons from giant-scale services.

In In IEEE Internet Computing, 2001.

[3] Fay Chang, Jeffrey Dean, Sanjay Ghemawat, Wilson C. Hsieh, Deborah A. Wallach, Mike Burrows, Tushar Chandra, Andrew Fikes, and Robert E. Gruber.

Bigtable: A distributed storage system for structured data.

In Proc. of USENIX OSDI, 2006.

[4] Jeffrey Dean and Sanjay Ghemawat.

Mapreduce: Simplified data processing on large clusters.

In Proc. of ACM OSDI, 2004.

References II

- [5] Lars George.
HBase, The Definitive Guide.
O'Reilly, 2011.
- [6] Sanjay Ghemawat, Howard Gobioff, and Shun-Tak Leung.
The google file system.
In *Proc. of ACM OSDI*, 2003.
- [7] Patrick O'Neil, Edward Cheng, Dieter Gawlick, and Elizabeth O'Neil.
The log-structured merge-tree (lsm-tree).
1996.

References III

[8] D. Salmen.

Cloud data structure diagramming techniques and design patterns.

<https://www.data-tactics-corp.com/index.php/component/jdownloads/finish/22-white-papers/68-cloud-data-structure-diagramming>, 2009.